

FinSight AI

Five-minute video walkthrough guide

Goal: show the reviewer what the project does, where the three assessment tasks live, and how the evidence was produced. Speak slowly and use your own words. Do not read private keys aloud.

Before you record

1. Open the folder `C:\Users\Admin\NSR\cdazz\finsight_ai` in your IDE.
2. Open the GitHub repository in a signed-out browser window:
`https://github.com/nipunarambukkanage/finsight_ai`.
3. Make sure the three notebooks show their outputs and that the generated artifacts are visible in each `artifacts` folder.
4. Open `README.md`, `CITATIONS.md`, and `REFLECTION.md` in separate tabs.
5. Close or hide any terminal that contains tokens, passwords, or Colab secrets.
6. Start a screen recorder. Select the IDE window and browser window. Keep the video under five minutes.

Important: show real outputs only. If a run is labelled demo or fixture, say so clearly. Never describe fixture output as live market data, teacher-generated data, or GPU training evidence.

Simple script and screen actions

Time	What to show	What to say
0:00-0:30	Root <code>README.md</code> and the repository name.	"This is FinSight AI. It is a modular financial research platform. The repository is split into three independent assessment tasks and a shared application core."
0:30-1:25	Open <code>task1_financial/notebooks/task1_equity_research.ipynb</code> and scroll through visible outputs. Show the artifact folder.	"Task 1 uses two years of market data, first-principles technical

		indicators, validated sentiment, and a one-page report. These outputs show the ticker, data provenance, indicator results, news reasoning, and the generated report.”
1:25-2:35	Open <code>task2_genai/notebooks/task2_filing_risk_qlora.ipynb</code> . Show the dataset split, loss chart or loss table, evaluation JSON, manual review, and model card.	“Task 2 extracts risks from filing excerpts. The student model is QLoRA fine-tuned with exact quotation grounding and abstention. This screen shows the teacher model record, source-disjoint data split, training losses, held-out metrics, and manual review.”
2:35-3:45	Open <code>task3_agentic/notebooks/task3_agentic_research.ipynb</code> . Show one tool call, its observation, the next action, the typed handoff, and the clarification response.	“Task 3 has a single researcher and a two-agent workflow. Data Analyst receives only data tools. Research Writer receives only news and web tools. The writer asks exactly one clarification question and uses the answer in the final report.”
3:45-4:25	Show <code>task3_agentic/logs/agent_trace.jsonl</code> , the cache folder, and the second-run result.	“Every tool call records the agent, validated arguments, a

		short output preview, duration, status, and trace ID. A repeated request uses the ticker and date cache and returns a cached result without repeating the tools.”
4:25-4:50	Open CITATIONS.md , REFLECTION.md , and the Task READMEs with Colab badges.	“The repository records AI assistance, external sources, the teacher prompt, design choices, and limitations. Each task README links directly to its Colab notebook.”
4:50-5:00	Show the GitHub page and stop the recording.	“The repository contains the code and the executed evidence. Thank you for reviewing FinSight AI.”

What the reviewer should see

Task	Minimum visible evidence
Task 1	Live data mode, two-year range, indicator table, at least ten headlines, validated sentiment, recommendation, provenance, and Markdown/HTML/chart artifacts.
Task 2	Teacher health check, 200 accepted examples, 80/10/10 source-disjoint split, QLoRA settings, train/validation losses, merged checkpoint, evaluation metrics, ten manual labels, hallucination rate, and model link.
Task 3	Observation-driven tool calls, failed-tool recovery if available, role restrictions, one clarification cycle, final report, cache reuse, and real JSONL trace entries.

After recording

1. Watch the video once. Check that your voice is clear and the total time is no more than five minutes.
2. Upload it as an unlisted YouTube video or to Google Drive with “Anyone with the link can view”.
3. Open the shared link in a signed-out browser window.
4. Copy the final video URL into the submission email or root README.

Final sentence to remember: “I will only claim evidence that is visible in the executed notebook or committed artifact.”

Quick final checklist

- ☐ Three notebooks executed with outputs visible
- ☐ Colab badges open correctly
- ☐ Task 1 report artifacts committed
- ☐ Task 2 dataset, losses, evaluation, manual review, merged model, and model card committed
- ☐ Task 3 real trace, report, handoff, clarification, and cache evidence committed
- ☐ GitHub, model, Drive, and video links tested while signed out
- ☐ No credentials appear on screen or in the repository

FinSight AI assessment walkthrough guide. Keep this guide with the repository under `docs/video-walkthrough-guide.pdf`.